



TaxiGreen

AI-Driven Fleet Optimization

Universidade do Minho | Grupo 06

João Delgado | Nelson Mendes | Simão Mendes | Tomás Machado

Presentation Agenda

01. Problem Description: System representation (S,A,T,C) and fleet management.

02. City Representation: Braga's graph, zonal classification, and critical infrastructure.

03. Search Strategies: Comparison between Uninformed and Informed search algorithms.

04. Dynamic Simulation: Real-time traffic, weather modeling, and random request generation.

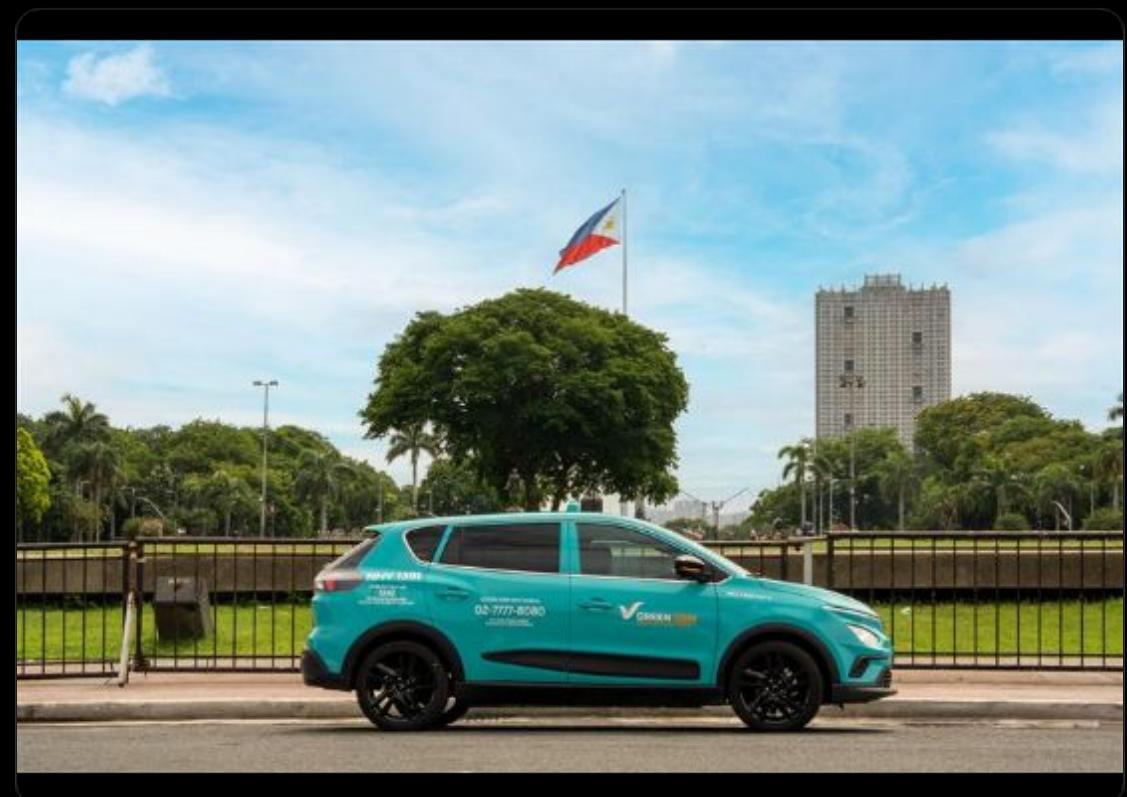
05. Performance Benchmark: Quantitative analysis of time, memory, and path optimality.

06. Architecture & Results: System design, multithreaded server, and final conclusions.

The Challenge: Sustainable Mobility

Goal: Optimize a mixed fleet (combustion & electric) in a real urban environment.

- ⚡ Manage distinct constraints: Range anxiety & charging times.
- 🏢 Operate within the complex topology of Braga, Portugal.
- ⚖️ Balance conflicting objectives: Minimizing time vs. Maximizing sustainability.



Problem Formalization (S , S_0 , A , T , C)



State (S)

Vehicle Location , Battery Level ,
Operational Status , Passenger
Capacity , Engine Type , Pending
Requests , Priority , Deadline ,
Environmental Preference , and
Dynamic Traffic Conditions.



Actions (A)

Move(v , dest), Pickup(v , p),
Dropoff(v , p), Refuel(v , s).



Goal (T)

Pending Requests Empty ($P = \emptyset$) ,
Successful Passenger Transport ,
Deadlines Met , and Guaranteed
Vehicle Autonomy.

Modeling Braga

We modeled the entire municipality using **OSMnx** and OpenStreetMap data.

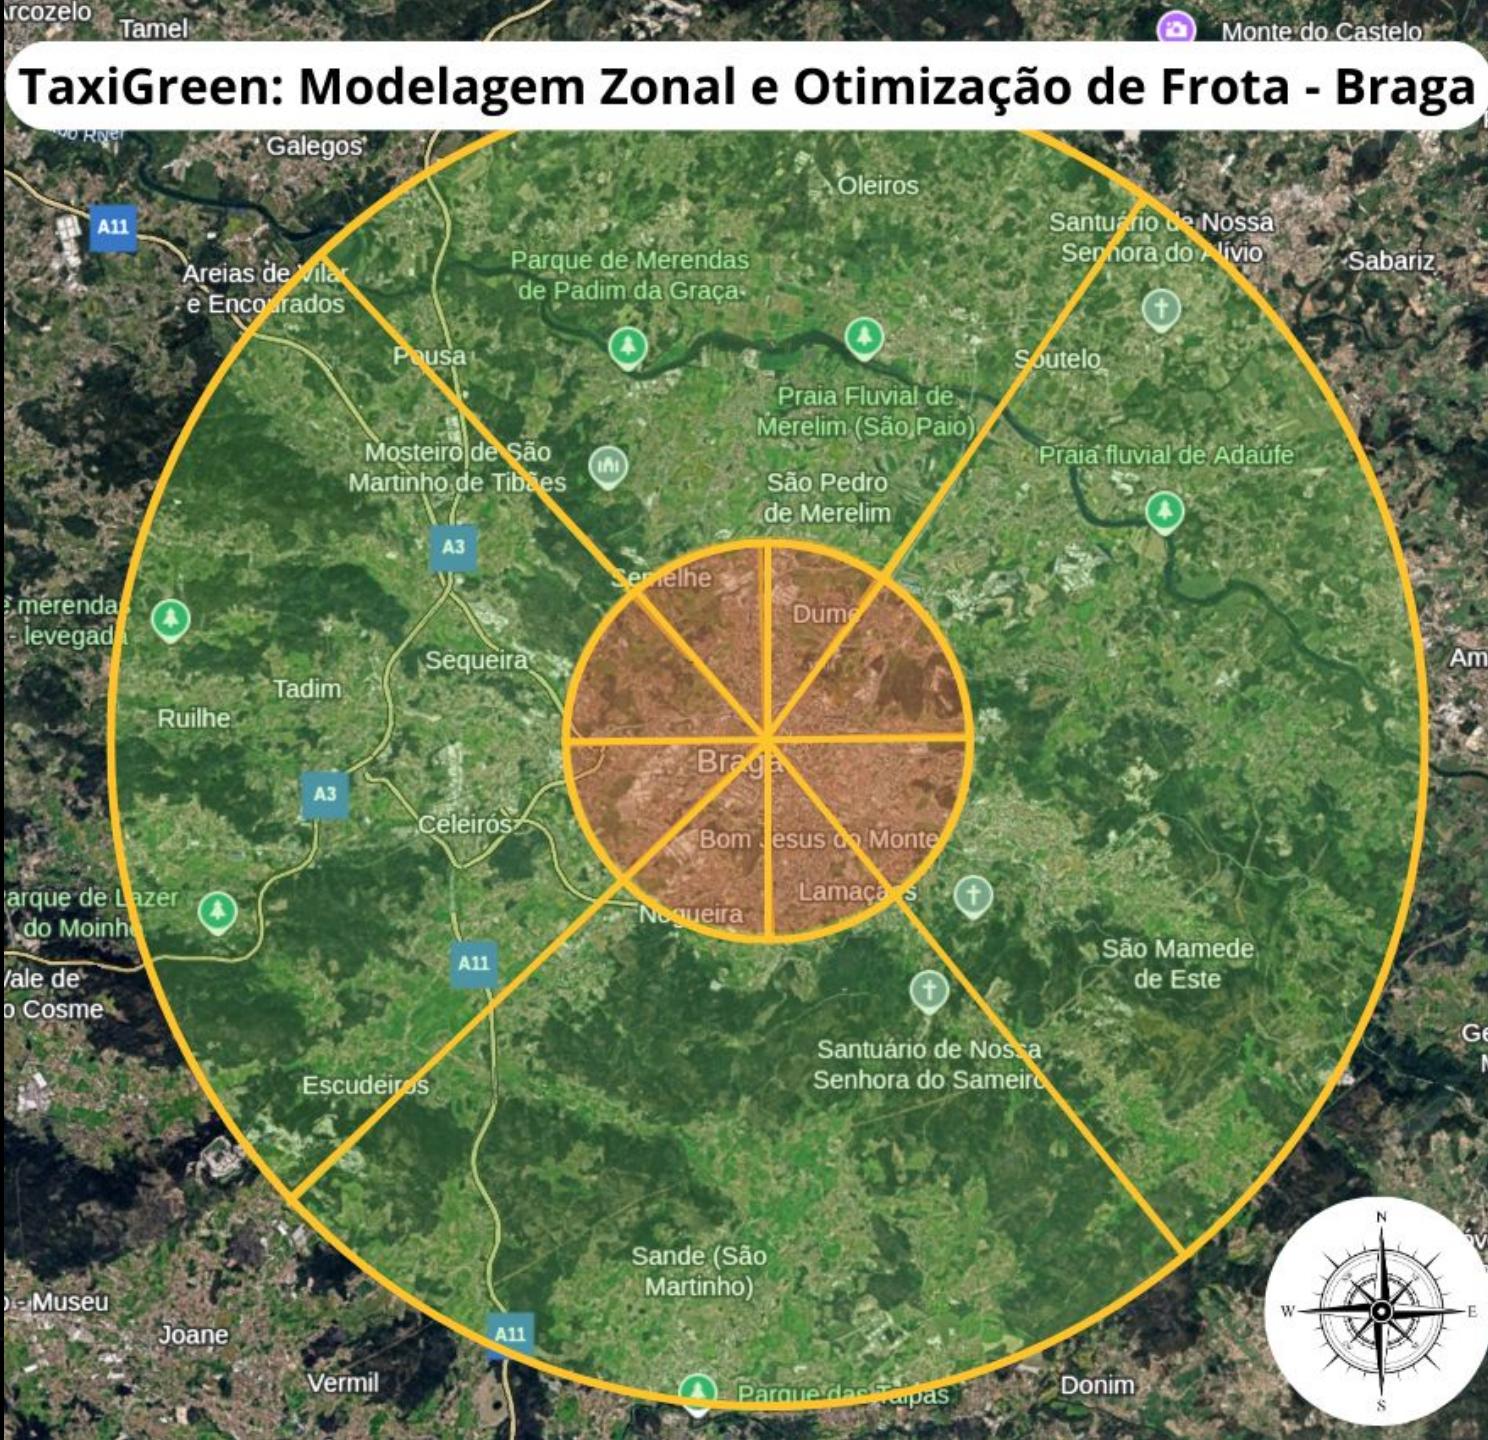
👉 **Graph:** 8,624 Nodes & 19,753 Edges.

📍 **Zonal Logic:**

Center: High density, 8 cardinal sectors (Traffic heavy).

Periphery: Rural, 4 sectors (Weather sensitive).

🌿 **Ecological Score:** dynamic (η) for each road segment.



The Eco-Score

Calculating Sustainability

Eco-Score Definition Dynamic Rating $\eta \in [0,1]$,

Environmental Weight ($\beta=0.2$), and Route Sustainability

 **High Score (0.9 - 1.0):** Residential streets, cycle lanes,
low-speed zones.

 **Medium Score (0.5 - 0.8):** Main avenues, secondary
distributors.

 **Low Score (0.0 - 0.4):** Highways, industrial zones,
high-speed arterials.



Effect on Routing

An edge with $\eta = 1$ incurs **zero penalty**.

An edge with $\eta = 0$ increases cost by **20%** (when $\beta = 0.2$).

Cost Function

Dynamic Cost Function

We calculate the cost of traversing an edge (u, v) by weighting time with environmental factors.

$$c(u, v) = [t_{\text{base}}(u, v) \times M_{\text{climate}}(u) \times M_{\text{traffic}}(u)] \times [1 + \beta \times (1 - \eta(u, v))]$$

Simplified Dynamic Cost Function

$$\text{Cost}_{\text{total}} = \text{Cost}_{\text{base}} \times M_{\text{climate}} \times M_{\text{traffic}} \times \text{Score}_{\text{Eco}}$$

Where:

- $t_{\text{base}}(u, v)$: represents the **minimum travel time** (in seconds), calculated by the ratio of the edge length to its maximum speed converted to m/s;
- M_{climate} e M_{traffic} : are **dynamic coefficients (≥ 1.0) determined by the geographical area** of the node of origin;
- $\beta = 0.2$: is the **weight given to the ecological component** of the system;
- $\eta(u, v) \in [0, 1]$: is the **Eco-Score** of the edge, where higher values indicate more environmentally friendly routes.

Heuristic

Admissible Heuristic

To ensure A* optimality, we use the Haversine distance divided by the global maximum speed.

$$h(u, v) = \frac{D_{\text{Haversine}}(u, v)}{\frac{v_{\text{global_max}}}{3.6}}$$

Where:

- $D_{\text{Haversine}}(u, v)$: Calculates the great-circle distance (straight line) between the current node and the destination.
- $v_{\text{global_max}}$: Represents the highest speed limit recorded in the entire graph (e.g., 120 km/h on highways).
- 3.6 Conversion Factor: Used to convert speed from km/h to m/s to ensure metric consistency.

Search Strategies

⌚ Uninformed Search

DFS: Explores deeply before backtracking. Gets lost in cycles or long paths. Result: $\approx +2000\%$ Cost.

BFS: Guarantees fewest hops, but ignores traffic/eco costs. Result: $\approx +31.8\%$ higher cost.

UCS: Guarantees optimal cost, but explores radially (slow $\approx 63.30\text{ms}$).

⌚ Informed Search

Greedy: Fastest ($\approx 3.05\text{ms}$), but suboptimal solutions ($\approx +27.5\%$ cost).

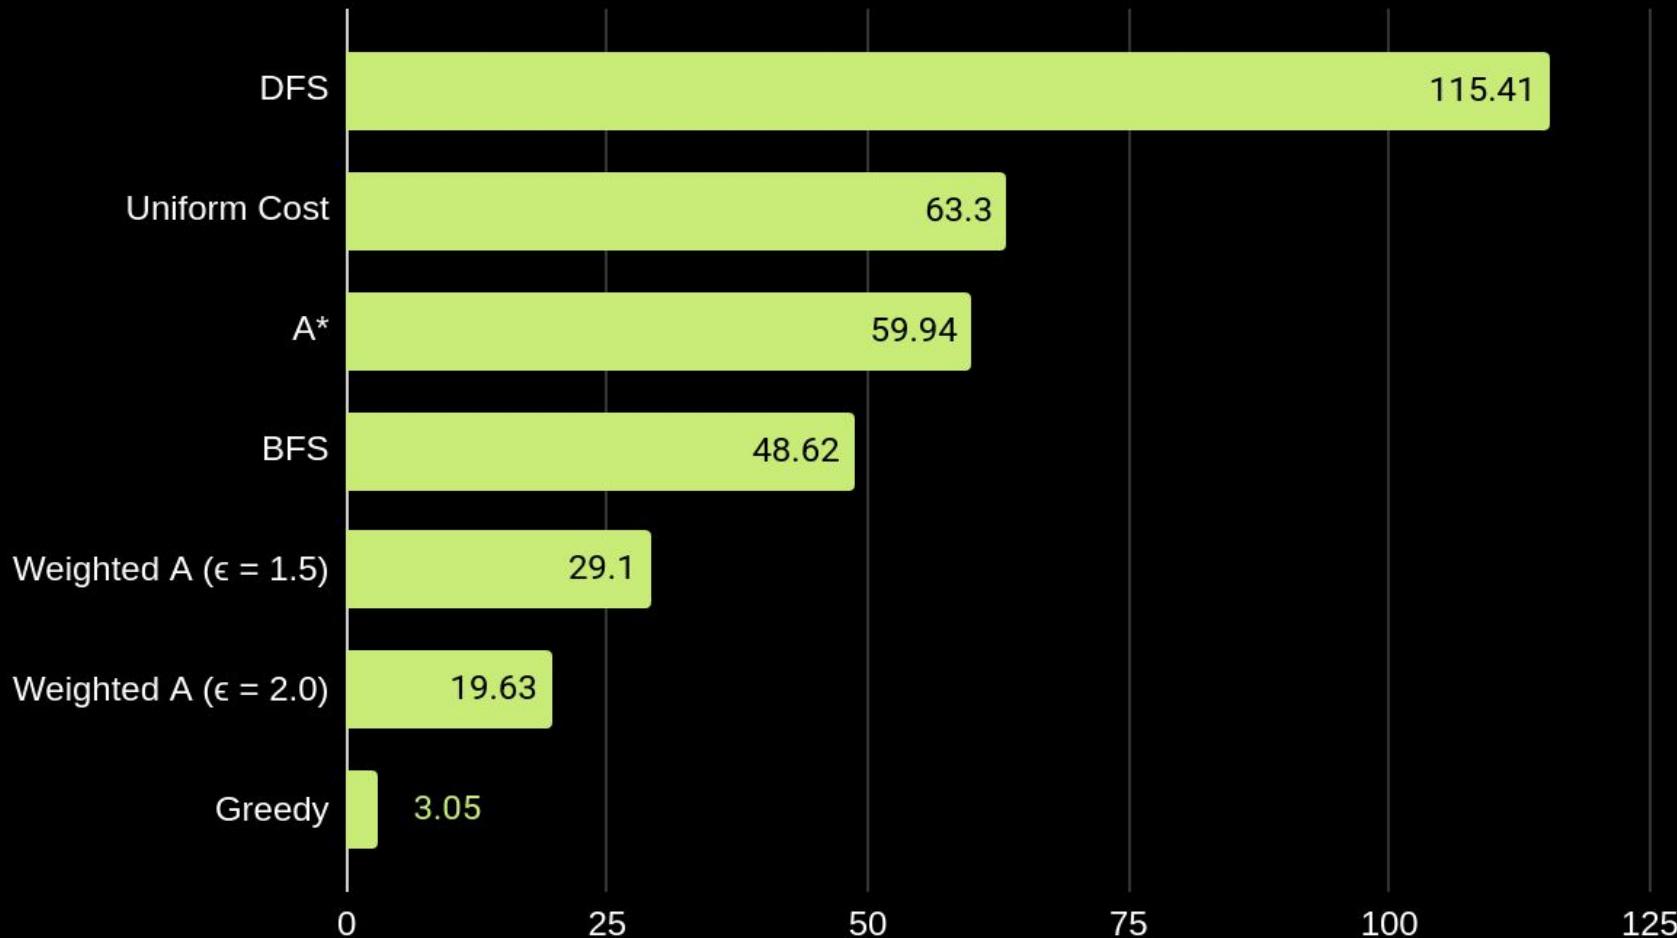
A* (A-Star): The gold standard. Combines UCS optimality with Greedy's direction.

Weighted A*: Prioritizes moving towards goal over minimizing path cost. **The Winner** ($\approx 19\text{ms}$ with $\epsilon = 2.0$).

Benchmarking Methodology

- ✓ **Isolated Environment:** Tests ran on a dedicated script, bypassing the UI/Network latency.
- ✓ **100 Iterations:** We generated 100 random (Origin, Destination) pairs.
- ✓ **Consistency:** Exactly 100 iterations per route to ensure results comparability.
- ✓ **Tools:** Python `time.perf_counter()` and `tracemalloc` for memory profiling.

Benchmarking: Execution Time (milliseconds)



Greedy is the fastest but not the best due to suboptimal results; however, A is the top performer, as shown in the next slide.*

The Trade-off: Speed vs. Optimality

Algoritmo	Sucesso	Tempo (ms)	Memória (KB)	Custo	Passos	Eficiência
A*	100%	59.54	504.5	559.2	65.9	Ótimo
Uniform Cost	100%	63.30	670.4	559.2	65.9	Ótimo
Weighted A* ($\epsilon = 1.5$)	100%	29.10	738.3	559.2	65.9	Ótimo
Weighted A* ($\epsilon = 2.0$)	100%	19.63	752.7	559.2	65.9	Ótimo
Greedy	100%	3.05	91.8	712.7	71.6	+27.5%
BFS	100%	48.62	418.2	737.1	44.1	+31.8%
DFS	100%	115.41	23,776.8	12,056.7	1314.0	+2056%

Optimization: Heuristic Caching

The computational cost of the Haversine formula initially slowed down A*. Caching solved this.

84ms

PRE-OPTIMIZATION

29.1%

PERFORMANCE BOOST

~59ms

POST-OPTIMIZATION

The Trade-off Analysis

The "Golden Mean"

We selected **Weighted A*** ($\epsilon = 2.0$) as our primary engine.

- **Speed:** 19.63ms (Extremely fast)
- **Cost:** 559.2 (Identical to Optimal)
- **Why?** In the topology of Braga, the slight "greediness" ($\epsilon = 2.0$) directs the search efficiently without getting trapped in suboptimal local minima. Weighted

3X

FASTER THAN A*

0%

COST PENALTY

System Architecture



Server

Multithreaded TCP Server (Port 8888).
Maintains simulation loop, by sending
a random request every 3 seconds.



Client

One-threaded client using socket and
pickle for serialization.
Decoupled from logic.



Shared State

Thread-safe locks protect the Fleet
and Request Queue during concurrent
access.

Rich Terminal UI

Developed using the **Python Rich** library, the interface offers a dashboard-like experience in the terminal.

- ✓ Live Fleet Status Tables
- ✓ Algorithm Switcher
- ✓ Real-time System Logs
- ✓ "Matrix" Aesthetic



Dynamic Simulation

The system is not static. It reacts to a chaotic environment to test fleet resilience.

Variable Traffic: Zonal multipliers simulate rush hours ($M_{\text{traffic}} > 2.0\$$).

Weather Events: Rain reduces average speed.

Infrastructure Failure: Charging stations or streets in general, can go down, forcing re-routing.



Conclusion

- ✓ **Modeling Success:** Graph representation with zonal logic effectively captures Braga's complexity.
- 🏆 **Algorithm Winner:** Weighted A* ($\epsilon=2.0$) offers the best balance of speed and optimality for real-time dispatch.
- 🌿 **Sustainability:** The dynamic cost function successfully prioritizes eco-friendly routes without compromising service levels.
- errupting / **Technical Robustness:** Threaded architecture ensures smooth operation under load.

Q&A

Thank you for your attention.

```
root@taxigreen:~$ shutdown -h now_
```



TaxiGreen

AI-Driven Fleet Optimization

Universidade do Minho | Grupo 06

João Delgado | Nelson Mendes | Simão Mendes | Tomás Machado